

## Lecture 4: January 15, 2018

*Lecturer: Lesley Istead**Notes By: Harsh Mistry*

## 4.1 Synchronization

### 4.1.1 Thread Synchronization

- All threads in a concurrent program share access to the programs global variables and the heap.
- The part of a concurrent program in which a shared object is accessed is called a critical section.
- To avoid conflicts known as race conditions we can enforce mutual exclusion and use variables
- Using `volatile` keyword forces the compiler to not optimize the code and forcefully load and store the value of the variable upon every use

### 4.1.2 Mutual Exclusion

Mutual exclusion is the process of defining a specific area of code that must execute without interruption. This can be achieved through locks.

#### Enforcing Mutual Exclusion With Locks

```
int volatile total = 0;
/* lock for total: false => free, true => locked */
bool volatile total_lock = false;

void add() {
    int i;
    for (i=0; i<N; i++) {
        Acquire(&total_lock);
        total++;
        Release(&total_lock);
    }
}

void sub() {
    int i;
    for (i=0; i<N; i++) {
        Acquire(&total_lock);
        total--;
        Release(&total_lock);
    }
}
```

---

Acquire/Release must ensure that only one thread at a time can hold the lock, even if both attempt to Acquire at the same time. If a thread cannot Acquire the lock immediately, it must wait until the lock is available.

---

Taken from lecture notes

Both Acquire and Release must be atomic functions, therefore, we can not rely on a software solution for efficient locking/releasing. As a result, we must use the CPU's hardware instruction to enforce mutual exclusion